

Big Data Analytics

Assignment 3

Due: Monday April 14th @ 11:59PM

Introduction

You will work with a large-scale dataset: the **McAuley-Lab/Amazon-Reviews-2023** dataset, hosted on Hugging Face at:

<https://huggingface.co/datasets/McAuley-Lab/Amazon-Reviews-2023>

The uncompressed size is roughly 200GB across 34 categories. You will complete a series of tasks spanning data ingestion, cleaning, exploratory data analysis (EDA), sentiment classification, recommender systems, and clustering.

All methods and algorithms are **fixed** (no substitutions), but you may implement these steps in any environment (local machine, cluster, or cloud). Work in groups of 3–4, dividing tasks as you see fit.

Dataset Fields:

For User Reviews (“review” data):

- **rating (float)**: User’s rating of the product (1.0–5.0)
- **title (str)**: Title of the user review
- **text (str)**: Text body of the user review
- **images (list)**: URLs or metadata of images in different sizes
- **asin (str)**: ID of the product
- **parent_asin (str)**: Parent ID of the product (products with different colors/sizes can share the same parent)
- **user_id (str)**: ID of the reviewer
- **timestamp (int)**: Time of the review (Unix time)
- **verified_purchase (bool)**: Whether the purchase was verified

- **helpful_vote (int)**: Number of helpful votes the review received
- **For Item Metadata** (“meta” data):
- **main_category (str)**: The domain or top-level category
- **title (str)**: Name of the product
- **average_rating (float)**: Rating shown on the product page
- **rating_number (int)**: Number of ratings for the product
- **features (list)**: Bullet-point features of the product
- **description (list)**: Descriptions of the product (often multi-line)
- **price (float)**: Price in US dollars (at time of crawling)
- **images (list)**: Product images (thumb, large, hi.res)
- **videos (list)**: Videos of the product (with title, URL)
- **store (str)**: Store name (where the product is sold)
- **categories (list)**: Hierarchical categories
- **details (dict)**: Product details (materials, brand, sizes, etc.)
- **parent_asin (str)**: Parent ID of the product
- **bought_together (list)**: Bundles or recommended items often bought together

Note: To link a user review with its product metadata, you typically use `parent_asin` as the join key. Some of the additional fields (like `helpful_vote`, `verified_purchase`, or `price`) can be incorporated into your analysis or EDA if you wish; however, the core tasks below are mandatory.

Assignment Specification

1. Data Acquisition

a) Obtain the Entire Dataset

You must download or otherwise access *all categories* (34) from **McAuley-Lab/Amazon-Reviews-2023**. The direct link is: <https://huggingface.co/datasets/McAuley-Lab/Amazon-Reviews-2023>. Use either:

- `load_dataset("McAuley-Lab/Amazon-Reviews-2023", ...)`
- The provided `download_all_amazon_reviews` function
- Or any other method, as long as you include *all categories*.

b) **Deliverable Proof**

Provide logs or screenshots verifying the successful acquisition of all categories (compressed or uncompressed).

c) **Deliverables (Data Acquisition):**

- *Report*: Evidence (screenshots, console output) showing you obtained all categories.
- *Code*: Scripts/notebooks demonstrating how you downloaded and organized the data.

2. Data Cleaning & Preprocessing

Perform these steps **exactly**:

a) **Merge on parent_asin**

Join review data with its matching metadata on `parent_asin` so each record combines user ID, product ID (`asin`), star rating, review text, verified purchase, helpful votes, brand/category info, etc.

b) **Handle Invalid / Missing Values**

- Drop rows where `star_rating` is missing or not in [1–5].
- Drop rows if `text` (the review body) is empty.
- If brand cannot be found in the metadata (e.g., missing in `details` or `store`), set `brand` = “Unknown”.

c) **Remove Duplicates**

If (`user_id`, `product_id`, `review_text`) repeats, keep only the first occurrence.

d) **Derived Columns:**

- **Review Length** = integer count of tokens in the `text` field (split on whitespace/punctuation).
- **Year** = extracted from `timestamp` if available (if there’s no timestamp, you may skip time-based EDA).

e) **Unified Output**

Consolidate all categories into a single large dataset (or partitioned store) for further tasks.

3. Exploratory Data Analysis (EDA)

Perform the following analyses on the *entire cleaned dataset* (all categories combined):

- a) **Star Rating Histogram:** Show a histogram for ratings 1–5.
- b) **Top 10 Categories:** Bar chart of categories by total review count (use the final `brand` or `main_category` if it is clearly identified).
- c) **Top 10 Brands:** Bar chart of brand by total review count (exclude “Unknown” from the top 10).
- d) **Time-Based Trend:** a line chart of average star rating per year.
- e) **Correlation:** Compute Pearson correlation between *review_length* and *star_rating*; state the numeric result and interpret briefly.
- f) **(Optional) Additional EDA Ideas:**
 - Distribution of `helpful_vote` counts.
 - Relationship between `verified_purchase` and star rating.

(These are not strictly required but may enrich your insights.)

4. Binary Sentiment Prediction (Logistic Regression)

Transform *rating* into:

Positive if rating > 3 , Negative if rating ≤ 3 .

- a) **Train/Test Split:** 80% train, 20% test, randomly shuffled.
- b) **Text Vectorization:** *TF-IDF* on review text (lowercase, split on whitespace/punctuation), discarding tokens in fewer than 5 reviews or in over 80% of reviews.
- c) **Classifier:** *Logistic Regression* (default hyperparameters).
- d) **Evaluation:**
 - Accuracy
 - F1 Score
 - Confusion Matrix (2×2: TP, FP, TN, FN)

5. Recommender System (ALS)

Create a **collaborative filtering** model using **Alternating Least Squares (ALS)**:

- a) **Data Setup**: Retain (user_id, product_id, rating). Drop users with fewer than 5 total reviews.
- b) **Split**: 80% train, 20% test.
- c) **ALS**: Use any library that supports ALS (Spark MLlib, etc.). Minimal tuning is acceptable.
- d) **Evaluation**: RMSE on the test set (predicted rating vs. actual).
- e) **Demo**: Show top 5 recommendations for 3 random users in the test set, including predicted ratings.

6. Clustering / Segmentation (k-means)

Segment *products* using **k-means** with $k = 5$:

- a) **Features** (per product):

(mean_rating, total_reviews, brand_id, category_id)

- **mean_rating**: Average user rating per product (*based on your merged data, not necessarily **average_rating** from metadata, though you could compare them.*)
 - **total_reviews**: Count of all reviews for that product
 - **brand_id**: Map each distinct brand string to an integer
 - **category_id**: Map each main category or top-level category string to an integer
- b) **k-means**: Exactly $k = 5$, default initialization, until convergence.
 - c) **Cluster Analysis**: For each cluster, report:
 - **Size**: number of products in the cluster
 - **Average mean_rating, average total_reviews**
 - **Average brand_id** and **category_id**
 - A short interpretation (e.g., high-rating electronics, unknown-brand items, etc.)

7. Final Report (8–15 pages)

Must have these sections in **exact** order:

- a) **Introduction**
- b) **Data Acquisition (Task 1)**
- c) **Data Cleaning & Preprocessing (Task 2)**
- d) **EDA (Task 3)**
- e) **Binary Sentiment (Logistic Regression) (Task 4)**
- f) **Recommender (ALS) (Task 5)**
- g) **Clustering (k-means) (Task 6)**
- h) **Conclusion (mention hardware/resources used, challenges, etc.)**

Code Repository: Organize notebooks/scripts clearly.

Grading Breakdown

Section	Weight
Data Acquisition (Task 1)	5%
Data Cleaning & Preprocessing (Task 2)	15%
EDA (Task 3)	10%
Binary Sentiment (Logistic Regression) (Task 4)	20%
Recommender (ALS) (Task 5)	20%
Clustering (k-means) (Task 6)	15%
Final Report (Task 7)	15%
Total	100%

Notes:

- **All Categories:** Ensure you eventually use all categories (~200GB total).
- **Resource Planning:** At least 16GB RAM is recommended and enough disk space or a streaming or distributed setup to handle large data.

Example Data Acquisition Workflows

These examples illustrate different ways to download/load data from McAuley-Lab/Amazon-Reviews-2023. Feel free to mix/modify to fit your environment.

Option A: Directly Use `load_dataset`

```
from datasets import load_dataset

# Example: Load "Automotive" reviews and metadata
auto_review = load_dataset(
    "McAuley-Lab/Amazon-Reviews-2023",
    "raw_review_Automotive",
    trust_remote_code=True
)
auto_meta = load_dataset(
    "McAuley-Lab/Amazon-Reviews-2023",
    "raw_meta_Automotive",
    trust_remote_code=True
)
```

Pros: Simple for smaller subsets or prototyping.

Cons: Repeated loads can be slow for the entire 200GB.

Option B: `save_to_disk` / `load_from_disk`

```
from datasets import load_dataset, load_from_disk
from pathlib import Path

# Download + save
auto_review = load_dataset(
    "McAuley-Lab/Amazon-Reviews-2023",
    "raw_review_Automotive",
    trust_remote_code=True
)
save_path = Path("/path/to/data/raw_review_Automotive")
auto_review.save_to_disk(str(save_path))

# Reload later
loaded_review = load_from_disk(str(save_path))
print(loaded_review)
```

Pros: One-time download, faster local reloads.

Cons: Requires storing all uncompressed data locally.

Option C: download_all_amazon_reviews (No Compression)

```
from bigdata_a3_utils import download_all_amazon_reviews

download_all_amazon_reviews(
    base_save_path="/path/to/data",
    categories=None,    # all categories
    compress=False
)
# Creates folders: raw_review_<CATEGORY>/, raw_meta_<CATEGORY>/, etc.
```

Then you can use `load_from_disk` on the created folders.

Option D: download_all_amazon_reviews (With Compression)

```
from bigdata_a3_utils import download_all_amazon_reviews, load_compressed_dataset

download_all_amazon_reviews(
    base_save_path="/path/to/data_compressed",
    categories=None,    # all
    compress=True,
    compression_format="gz", # or "bz2", "xz"
    compression_level=6
)

# After downloading, you'll have .tar.gz files like raw_review_Automotive.tar.gz
dataset = load_compressed_dataset("/path/to/data_compressed/raw_review_Automotive.tar.gz")
```

Pros: Saves disk space.

Cons: More CPU/time overhead for compress/decompress.

Note: If using options **A** or **B**, `load_dataset()`, data is stored in your HuggingFace cache directory, so you do not need to redownload it each time you call `load_dataset()`. However, this directory is not cleared automatically. Use functions from `bigdata_a3_utils` to locate and/or clear your HuggingFace cache when you are finished with the data.